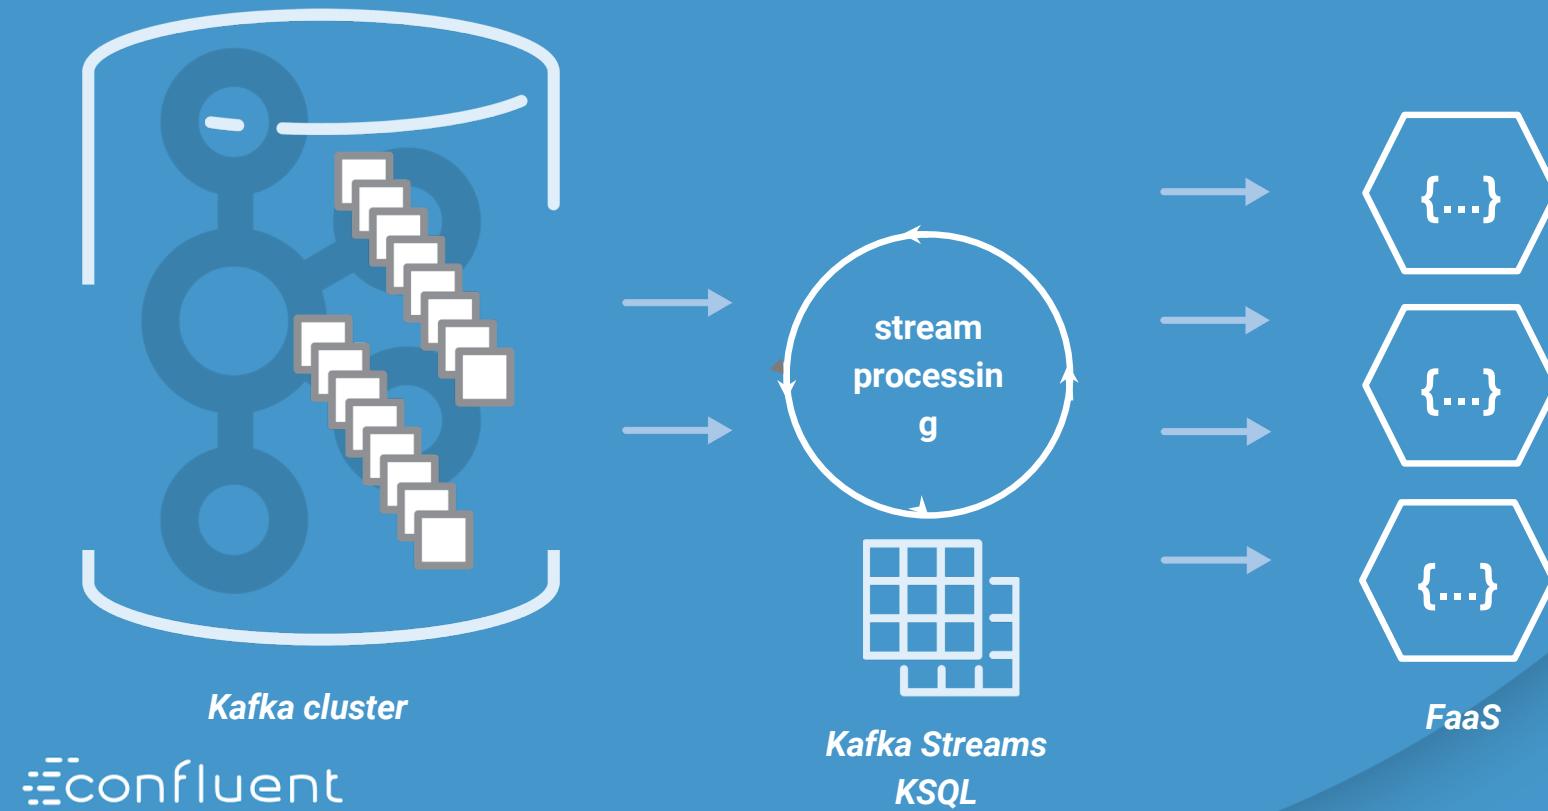




Event driven architectures with Kafka

A stack view

from stateful to stateless



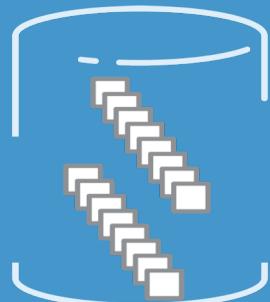
confluent

Stream processors versus FaaS processors

Kafka Streams & KSQL

- Stateful (and stateless)
- In band
- Dataflow - control plane
- Not ad hoc (except KSQL)
- Elastic

Think: data-model -->



confluent

FaaS - GCP Fn, Azure Fn, AWS Lambda etc

- Stateless
- Out of band (generally)
- Edge (in band)
- Ad hoc
- Super-elastic

Think: fire and forget -->



Events

What is an event?

**SOMETHING
HAPPENED!**

FACT!



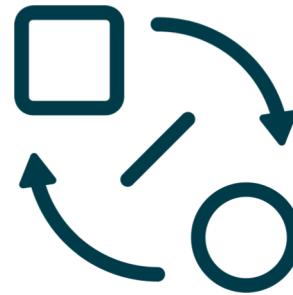
Events



A Sale



An Invoice



A Trade



A Customer
Experience

Events

Why do you care?

Loose coupling, autonomy, evolvability, scalability, resilience, traceability, replayability

...more importantly...

**BEING EVENT-FIRST CHANGES HOW YOU
THINK ABOUT WHAT YOU ARE BUILDING**



Events versus commands

Events:

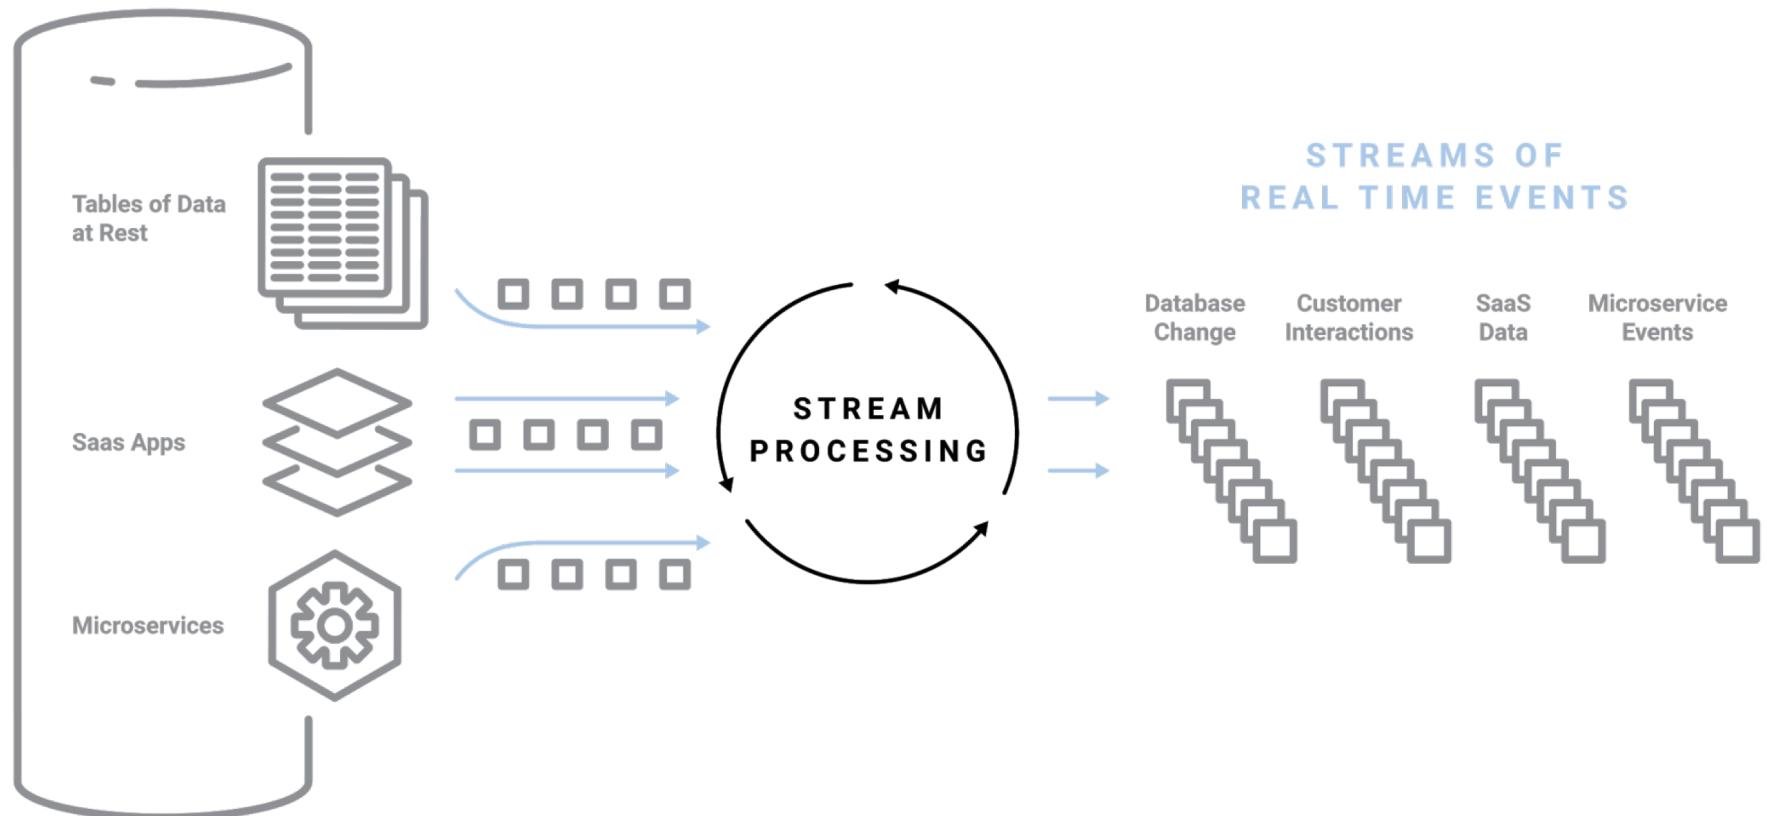
- Fact
- Something happen
- Immutable
- When
- Behavioural

Commands:

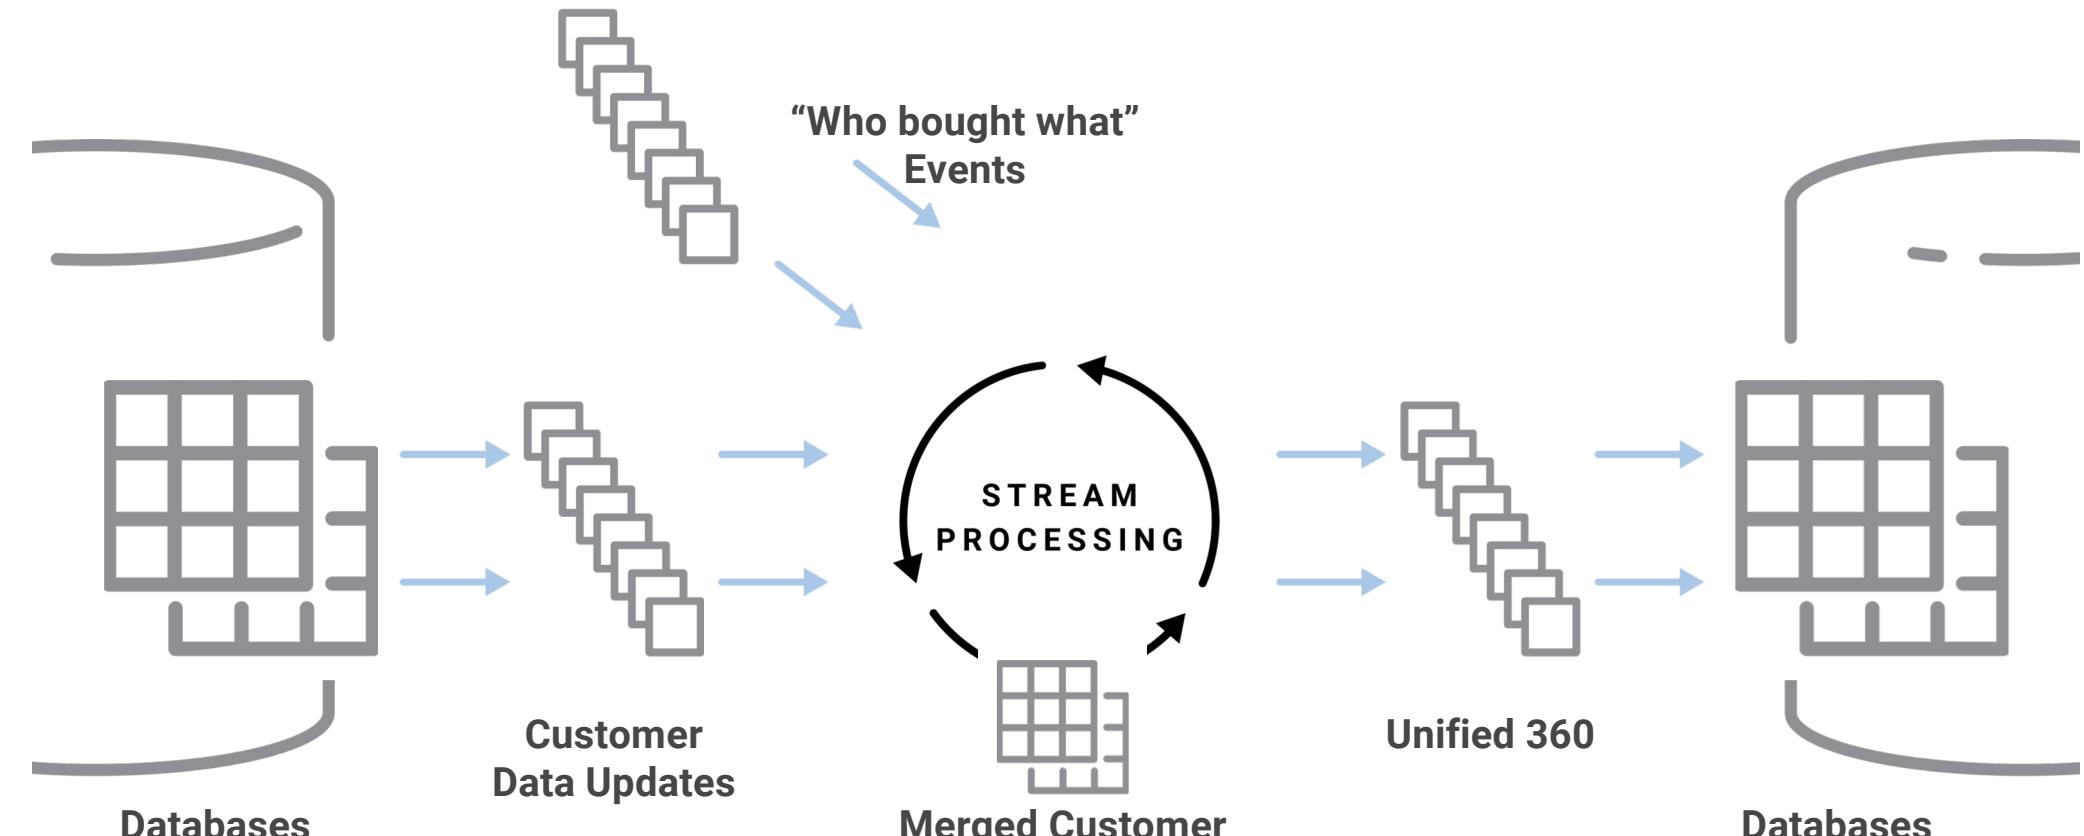
- Intent
- Contract
- Do something
- Coupling
- Structural

**EVENTS ARE USED TO MODEL BEHAVIOUR IN A DOMAIN
EVENT-SOURCING CAPTURES THAT BEHAVIOUR**

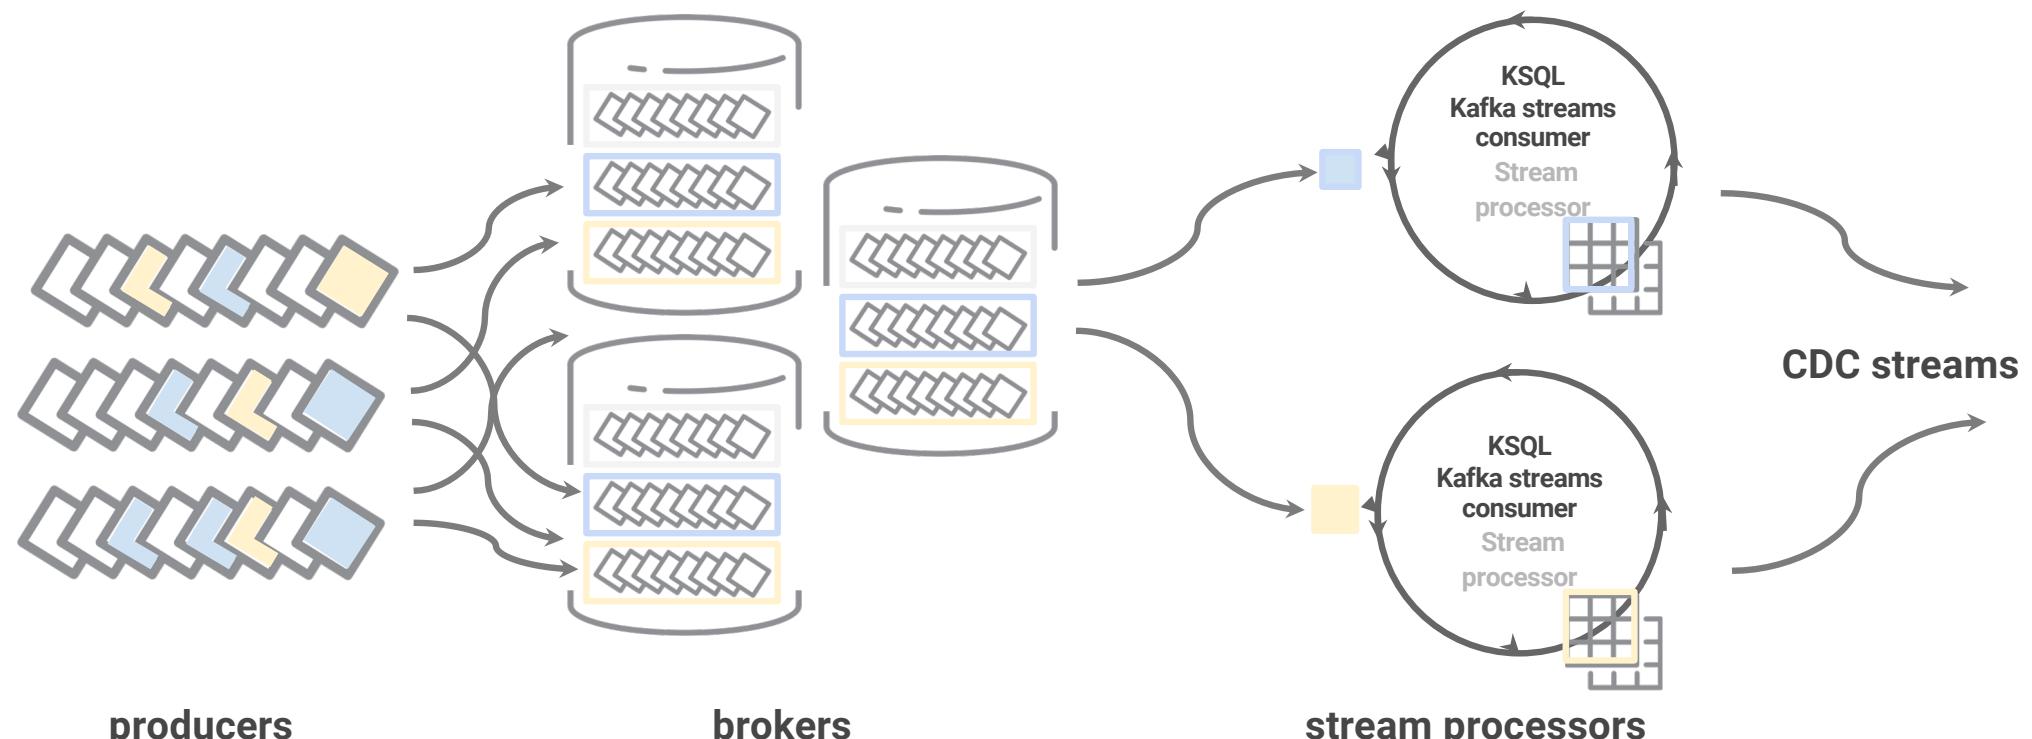
All your data is a stream of events



Events and data pipelines: Events as data



Events, Streams, Partitions, Tables

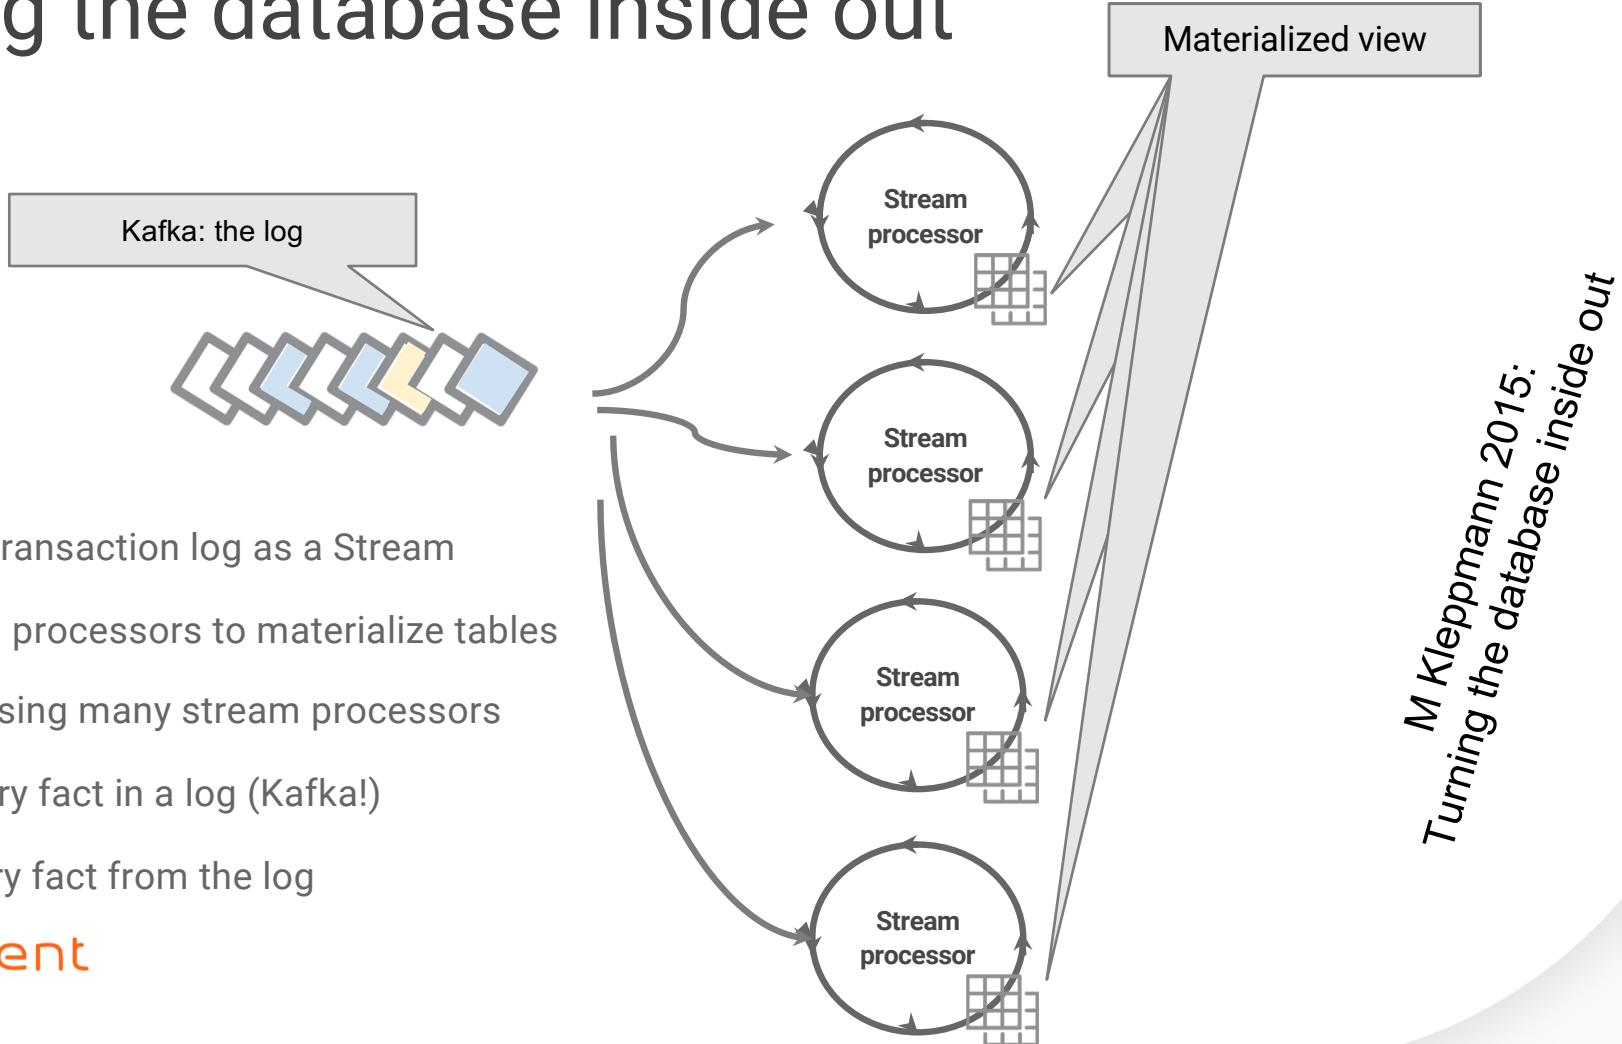


confluent

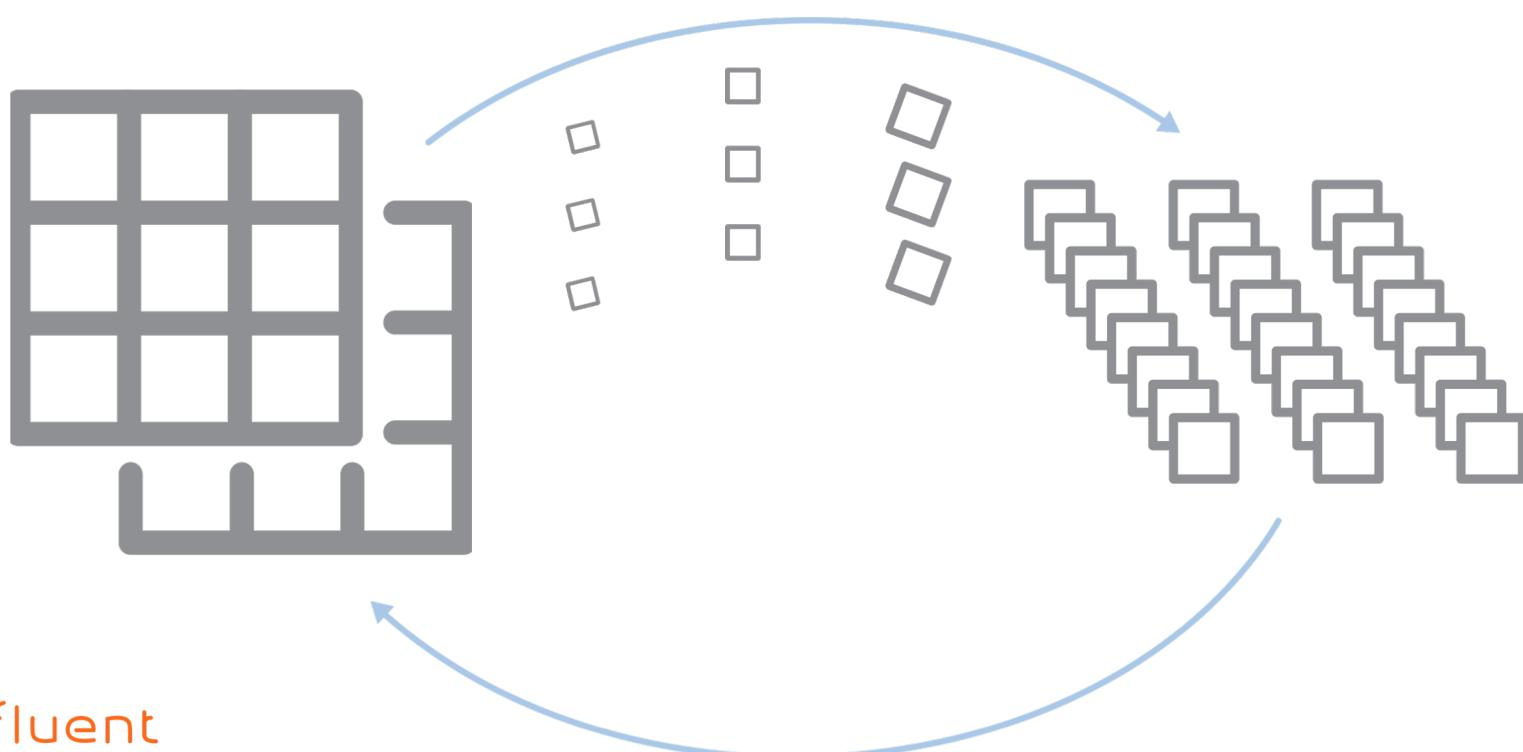
Turning the database inside out

- Model the transaction log as a Stream
- Use stream processors to materialize tables
- Scale out using many stream processors
- Record every fact in a log (Kafka!)
- Replay every fact from the log

confluent

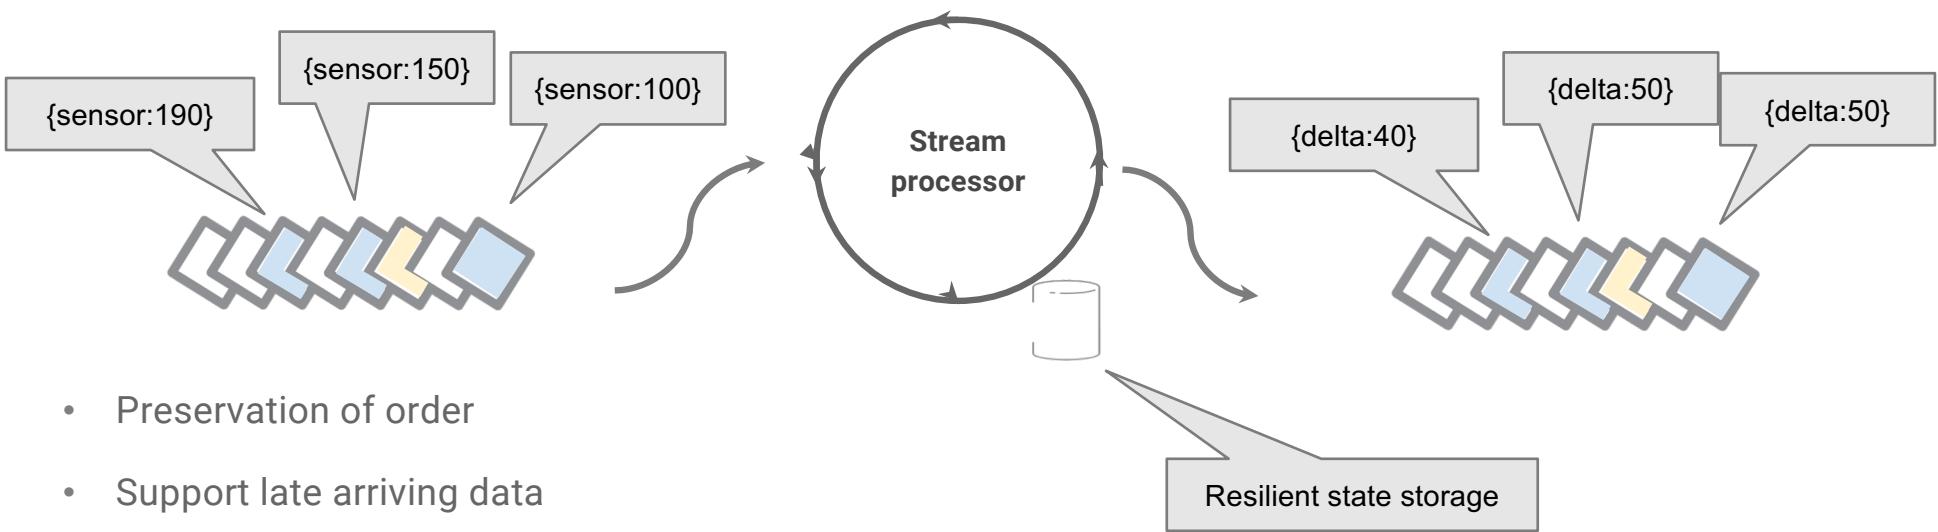


You need both tables and streams



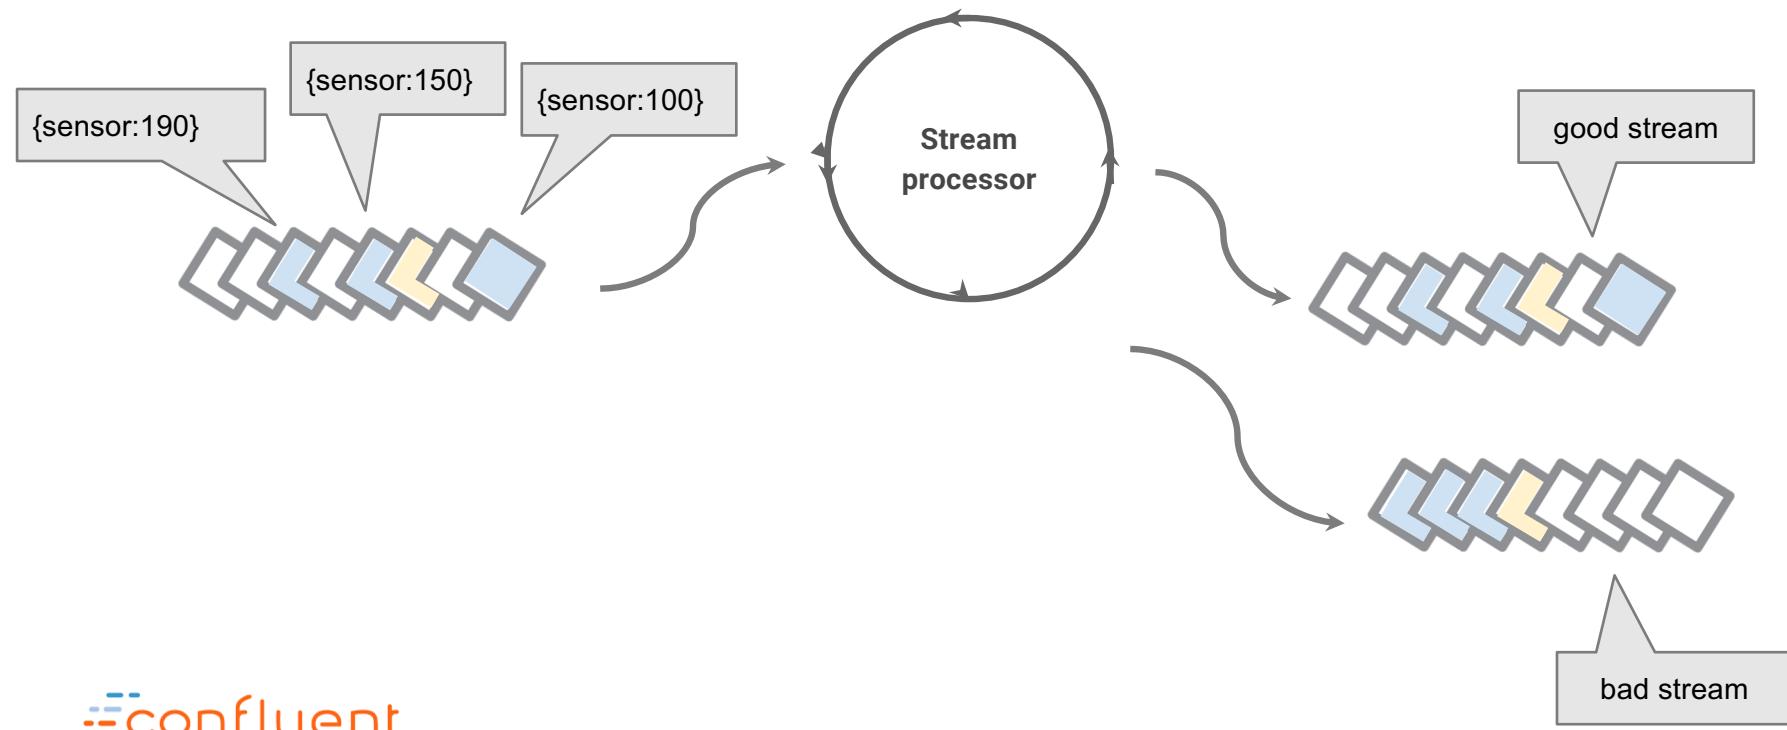
confluent

Streaming correctness (stateful)



- Preservation of order
- Support late arriving data
- Tolerance for out-of-sequence data
- Support machine failure scenarios
- Exactly once --- plus the ability to reason about time

Streaming correctness (stateless)

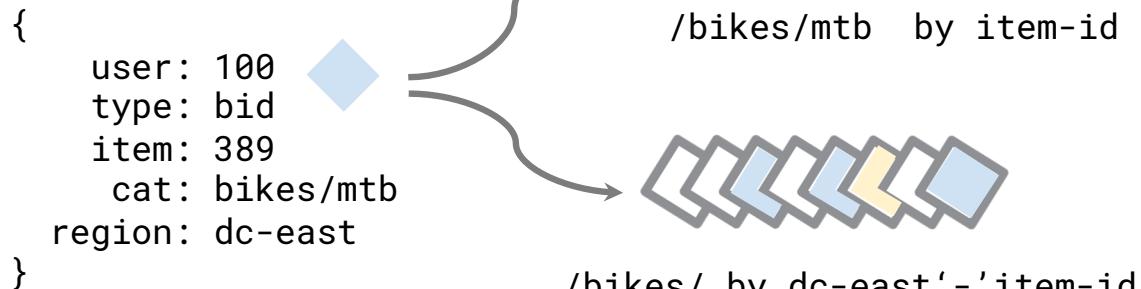


Data Modelling

Think events not commands,

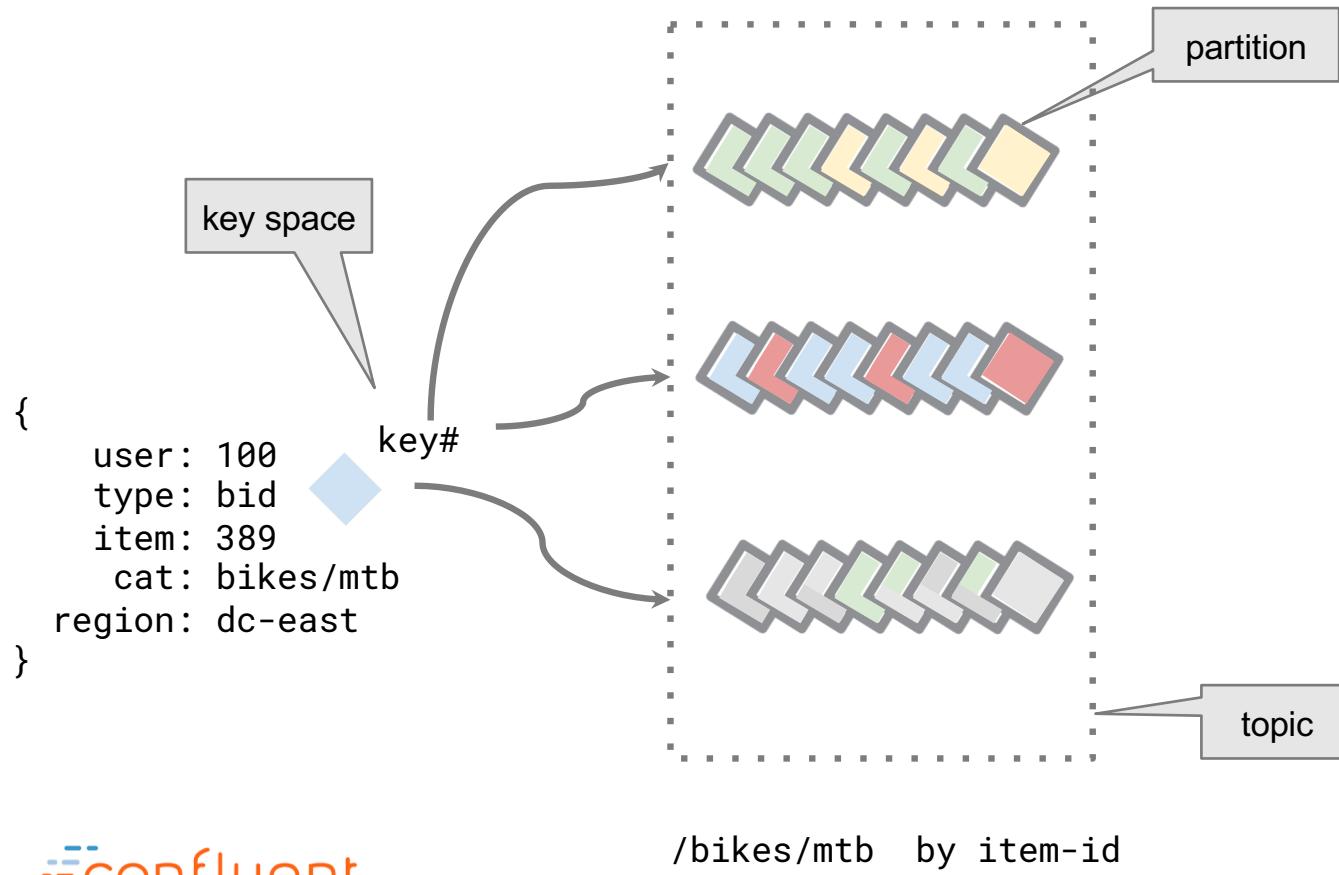
-> streams of events,

-> series of streams to model the domain



- Keyspace
- Throughput (events per sec)
- Throughput (historic)
- Parallelism
- Replicas
- Retention
- Data evolution

Data Modelling



- Keyspace
- Throughput (events per sec)
- Throughput (historic)
- Parallelism
- Replicas
- Retention
- Data evolution

Stateless versus Stateful

Stateless

- Filtering
- Transform
- Projection

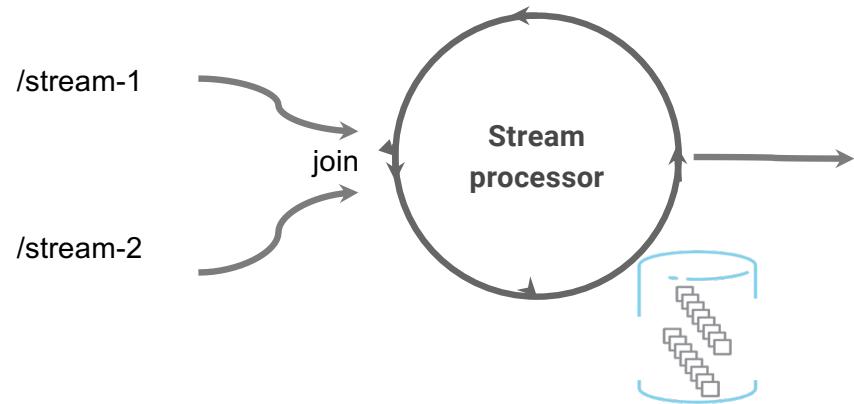
Think: fire and forget!



Stateful

- Window
- Aggregate
- Join (stream-table, stream-stream etc)

Think: remembers events!

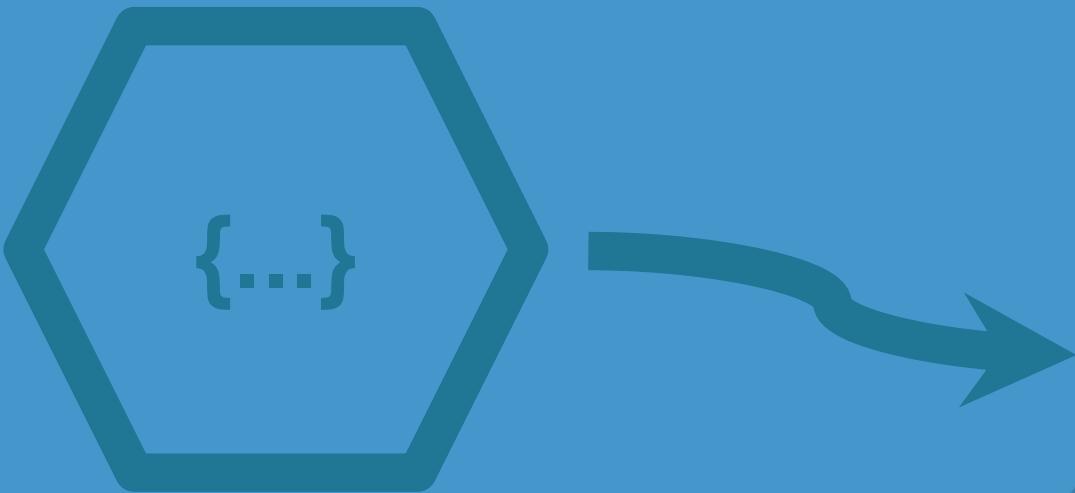


Winning with Event-driven streaming architectures

the secret sauce

the log, event sourcing, source of truth,
CQRS, event collaboration, replayability,
at-least once, exactly one, evolutionary
architectures, data-virtualization and
more

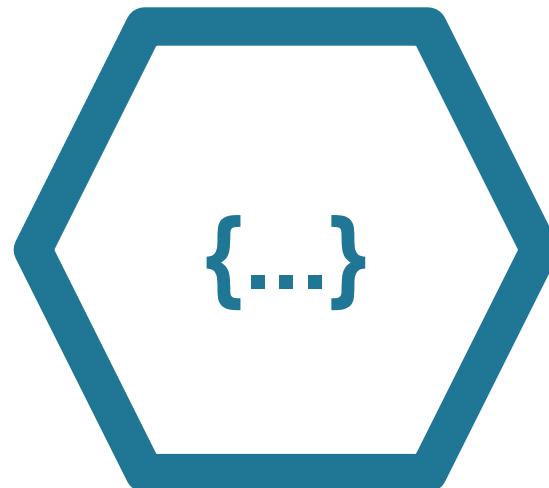




FaaS and Stream processing

What is FaaS?

- Bring your own compute
- Elastic
- Pay per use
- Stateless
- Cloud native



Any language, concurrent, sync or async



FaaS apps

1. Single site, events in any order

Async chained, **stateless**

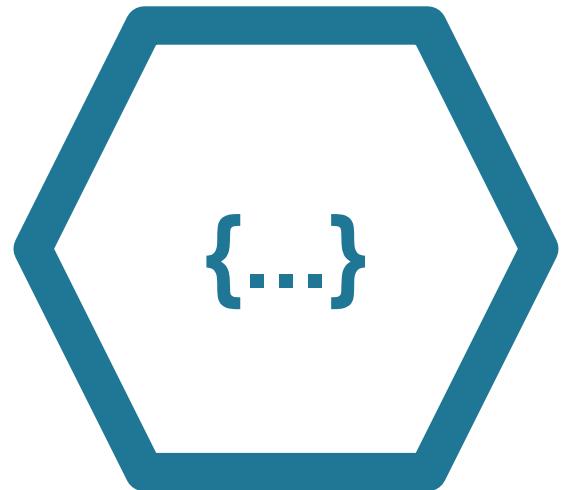
2. Non-stream oriented, **non-time critical**

3. Edge processing (in or out)

4. Enrich **incoming** events (stateless etl)

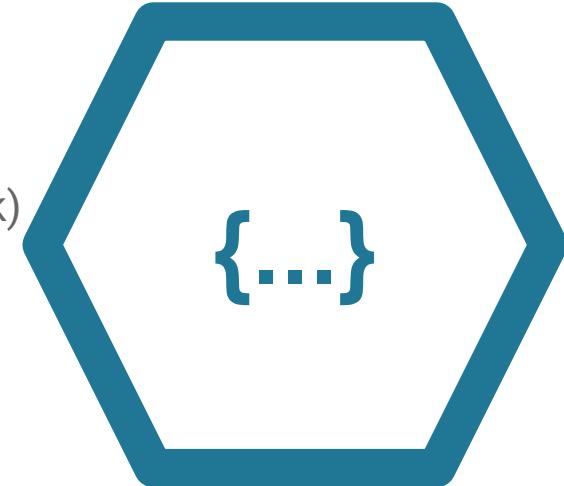
or

outgoing - email users, enrich image, ad-hoc requests



FaaS BUTs

- Sync chaining FaaS anti-pattern
- Complicated on different platforms
- Really really granular - like micronano-services
- Not cross cloud interoperable yet (without kafka or svrless fwk)
- Testability sucks
- Automation sucks



... but ...

FaaS requires the event-driven paradigm in order to be successful.

You cannot throw away 25 years of event-driven legacy and go back to rpc (fail!)

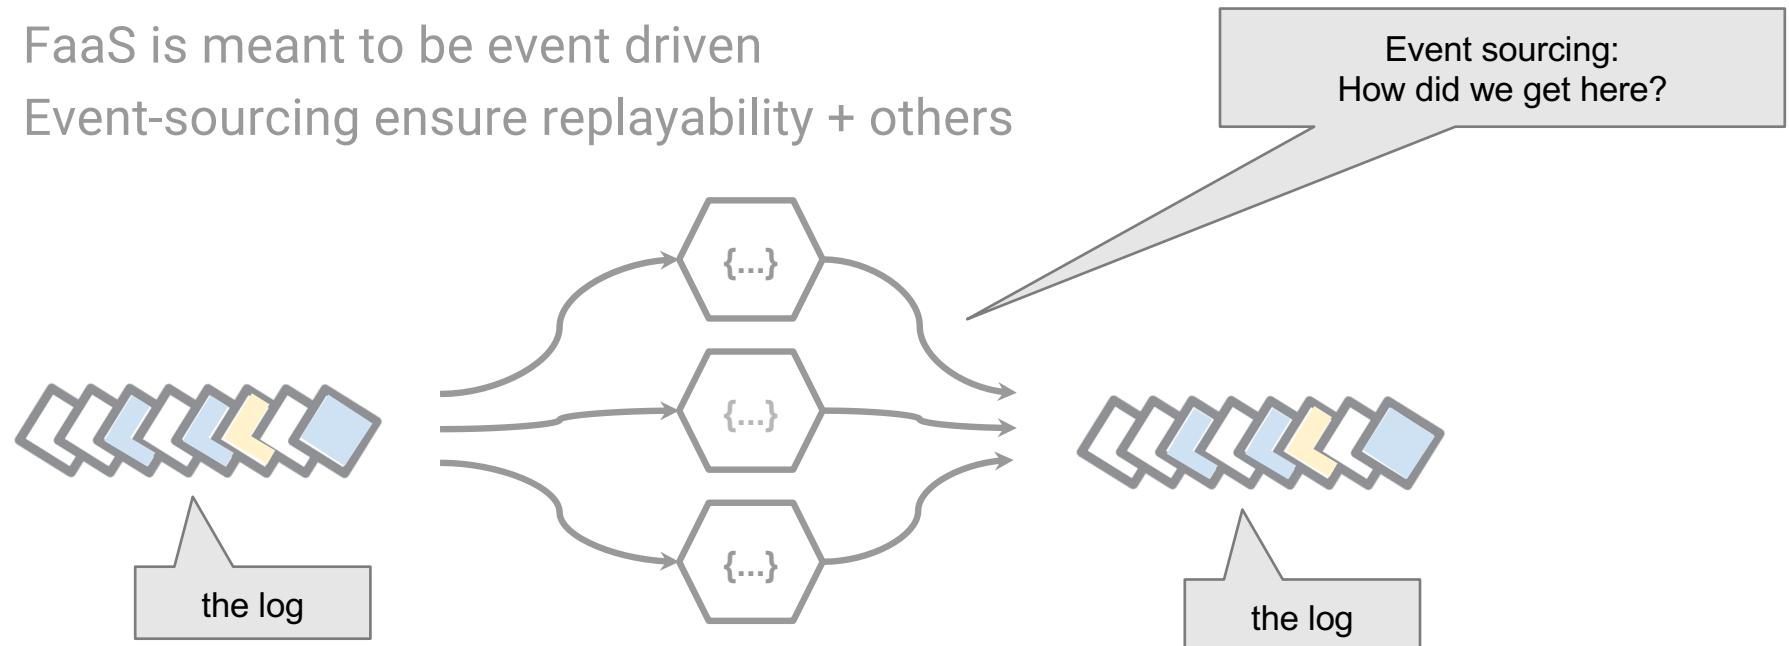
It's not 'event driven' unless you use event-sourcing



FaaS and event-sourcing

Recap:

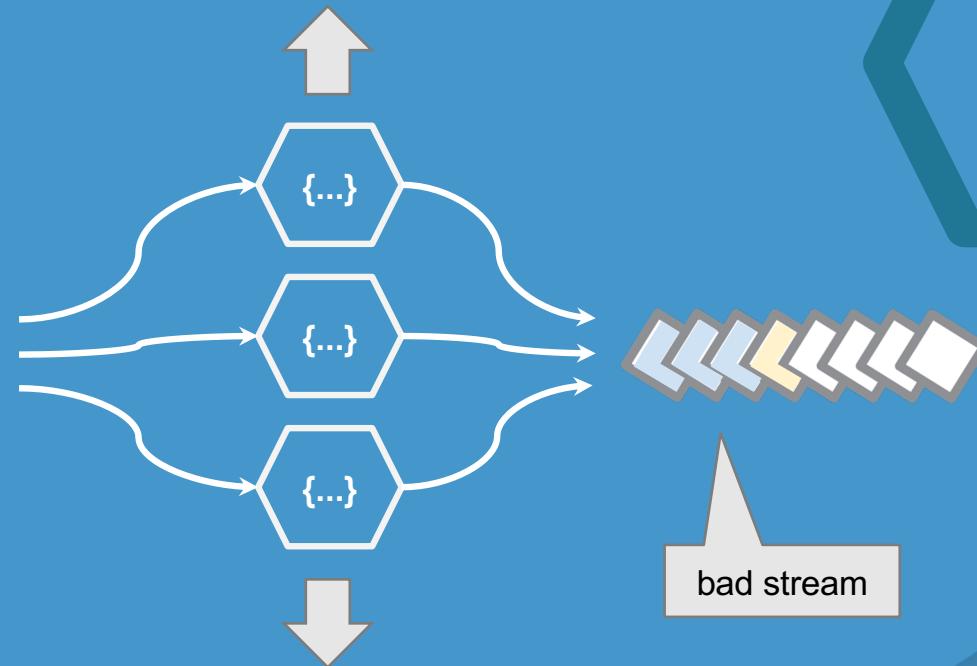
- Events are behavioural
- FaaS is meant to be event driven
- Event-sourcing ensure replayability + others



FaaS and streaming correctness

FaaS qualities

- Elastic
- Concurrency
- Sync or async
- Short-lived



FaaS for Stream processing

Needs...

- throughput
- ordering
- concurrency
- async

Works when...

- throughput per stream > invocation
- **stateless**
- not-historic processing
- async

Otherwise...

- pin streams to individual faas processors
- only 1 faas per stream
- problematic for scale :(



FaaS and the auction platform

- Enrich users on signup (address validation)
 - Geo-enrich items on placement (city, state, lat-lon cell identification)
 - Notify users on item sold or reserve not met
 - Perform analytics on auction when item ‘completes’
-
- Notify user of items-of-interest from their history when browsing
 - Ad-placement analytics (watched items, interested items, users-purchased)
-
- Monte-carlo auction simulations to guide users and calc item trending scores

In-band but edge

ad-hoc

Out of band, edge

Patterns for Infrastructure

1. **Ops:** (Observability, instrumentation, metrification) = monitoring patterns, dead-letter-queues, error-queues, audit-logs, application-logging, data-lineage
2. **Stream-based:** Worker Queue (compute grid or faas), event-backbone
3. **Data-based:** K/V store, Queryable Data fabric,
Data virtualization (via connect cdc streams) → Data Fabric
4. **FaaS:** fire and forget, event sourcing, unit-of-work



Key takeaways:

- **Model:** Events as the API, model the use-cases, model for scale, evolve the data-model : DDD
- **As Streams:** Streams are the database, tables materialized views, architect for evolutionary apps by using events
- **App Infra:** Build patterns to underpin higher order models (metrics, k/v etc)
- **FaaS:** Ad-hoc and edge stream processing or pinned processors

